

▼ Assignment 4

Report

```
1  %matplotlib inline
2  %config InlineBackend.figure_format = 'retina'
3
4  import numpy as np
5  import torch
6  from torch import nn
7  import torch.nn.functional as F
8  import pandas as pd
9  from sklearn.preprocessing import StandardScaler
10 import matplotlib.pyplot as plt
```

```
1  %tensorflow_version 2.x
2  from __future__ import absolute_import, division, print_function, unicode_literals
3  import functools
4  import os
5  import pandas as pd
6  import matplotlib.pyplot as plt
7  from keras.datasets import mnist
8  from matplotlib import pyplot
9
10 import numpy as np
11 import tensorflow as tf
12 import pandas as pd
13 import numpy as np
14 import torch
15 import torchvision
16 import torch.nn as nn
17 import torch.nn.functional as F
18 import torch.optim as optim
19
20 from torchvision.transforms import transforms
21 from torch.utils.data import DataLoader
22 from torch.utils.data import Dataset
```

```
1  %matplotlib inline
2  %config InlineBackend.figure_format = 'retina'
3
4  import numpy as np
5  import torch
6  from torch import nn
7  import torch.nn.functional as F
8  import pandas as pd
9  from sklearn.preprocessing import StandardScaler
10 import matplotlib.pyplot as plt
11 # Reading the data and selecting the rows
12 data = pd.read_csv('17078.csv')
13 x = data.iloc[:,0:784].values
14 y = data.iloc[:,784].values
15 print(len(y))
16
```

```

16
17 #Splitting the data in training and test data
18 from sklearn.model_selection import train_test_split
19 x_train,x_test,y_train,y_test = train_test_split(x,y, test_size = 0.25, random_state = 0,shuffle=True)
20
21
22

```


3000

Visualizing the data

```

1
2 pixels = x_train[9].reshape(28,28)
3 plt.imshow(pixels)
4 print(y_train[9])

```



Preprocessing the data

Scaling these values to a range of 0 to 1 before feeding them to the neural network model. To do so, I divided the values by 255

```

1
2 x_train = x_train/255
3 x_test = x_test/255

```

```

1 pixels = x_train[1].reshape(28,28)
2 plt.imshow(pixels)

```

<matplotlib.image.AxesImage at 0x7f3bb648d630>



▼ Building the model

I am building a fully connected neural network with input dimension 784 and output dimension is 10. It only has one hidden layer

The input and hidden layer has sigmoid activation function because the train data contain both positive and negative values.

The output layer has Softmax activation function.

```
1 import keras
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense
4 from torch import nn
5 bclfr = Sequential()
6
7 bclfr.add(Dense(units = 128 ,kernel_initializer = 'uniform',activation = 'sigmoid' , input_dim =784))
8 bclfr.add(Dense(units = 64 , kernel_initializer = 'uniform',activation = 'sigmoid'))
9 bclfr.add(Dense(units = 10 ,kernel_initializer='uniform',activation='softmax'))
```

I am using stochastic gradient descent for optimizing with learning rate of 0.01 and sparse_categorical_crossentropy as a loss function

```
1 bclfr.compile(optimizer=tf.keras.optimizers.Adam(0.01),loss='sparse_categorical_crossentropy',metrics=['accuracy'])
2 history = []
3 history.append(bclfr.fit(x_train,y_train,batch_size=64,epochs = 10,validation_data=(x_test,y_test))):
```



```
Epoch 1/10
36/36 [=====] - 0s 6ms/step - loss: 2.0368 - accuracy: 0.2364 - val_loss: 1.3175 - val_accuracy: 0.2364
Epoch 2/10
36/36 [=====] - 0s 4ms/step - loss: 0.8647 - accuracy: 0.7453 - val_loss: 0.6111 - val_accuracy: 0.7453
Epoch 3/10
36/36 [=====] - 0s 3ms/step - loss: 0.4240 - accuracy: 0.8813 - val_loss: 0.3982 - val_accuracy: 0.8813
Epoch 4/10
36/36 [=====] - 0s 4ms/step - loss: 0.2579 - accuracy: 0.9284 - val_loss: 0.3858 - val_accuracy: 0.9284
Epoch 5/10
36/36 [=====] - 0s 4ms/step - loss: 0.1641 - accuracy: 0.9551 - val_loss: 0.3916 - val_accuracy: 0.9551
Epoch 6/10
36/36 [=====] - 0s 3ms/step - loss: 0.1165 - accuracy: 0.9720 - val_loss: 0.3515 - val_accuracy: 0.9720
Epoch 7/10
36/36 [=====] - 0s 4ms/step - loss: 0.0772 - accuracy: 0.9813 - val_loss: 0.3068 - val_accuracy: 0.9813
Epoch 8/10
36/36 [=====] - 0s 4ms/step - loss: 0.0389 - accuracy: 0.9933 - val_loss: 0.3104 - val_accuracy: 0.9933
Epoch 9/10
36/36 [=====] - 0s 4ms/step - loss: 0.0253 - accuracy: 0.9973 - val_loss: 0.3033 - val_accuracy: 0.9973
Epoch 10/10
36/36 [=====] - 0s 3ms/step - loss: 0.0166 - accuracy: 0.9982 - val_loss: 0.3202 - val_accuracy: 0.9982
```

```
1 # plot diagnostic learning curves
2 def summarize_diagnostics(histories):
3     for i in range(len(histories)):
4         # plot loss
```

```


5  pyplot.subplot(2, 1, 1)
6  pyplot.title('Cross Entropy Loss')
7  pyplot.plot(histories[i].history['loss'], color='blue', label='train')
8  pyplot.plot(histories[i].history['val_loss'], color='orange', label='test')
9  pyplot.legend()
10 # plot accuracy
11 print('      ')
12 pyplot.subplot(2, 1, 2)
13 pyplot.title('Classification Accuracy')
14 pyplot.plot(histories[i].history['accuracy'], color='blue', label='train')
15 pyplot.plot(histories[i].history['val_accuracy'], color='orange', label='test')
16 pyplot.legend()
17
18 pyplot.show()
19

```

```

1  summarize_diagnostics(history)

```



▼ Observation

Form the above graph and above output of training we can see that at epoch 6 the validation accuracy is 91% which is highest of all , after epoch 6 the accuracy on the training set increases but the validation accuracy starts to decrease which is sign of **Overfitting**.

At epoch 1 or 2 we can see that both of the training accuracy and validation accuracy is two low which shows underfitting.

Even from the graph, we can easily see the underfitting of the model.

```

1  test_loss,test_acc = bclfr.evaluate(x_test,y_test,verbose=2)

```



```

24/24 - 0s - loss: 0.3646 - accuracy: 0.9107

```

```

1  probability_model = tf.keras.Sequential([bclfr])
2  predictions = probability_model.predict(x_test)

```

This function will plot a graph based on the results from the model trained above. It takes

- i - position
- prediction_array - The array of prediction predicted by my model.
- true_label - Actual label of test data set.

If the model predicted wrong result then the graph will be red, if the predicted result is right then the graph will be blue, if the model predicted some other value but not confident about it then the graph is in the grey color.

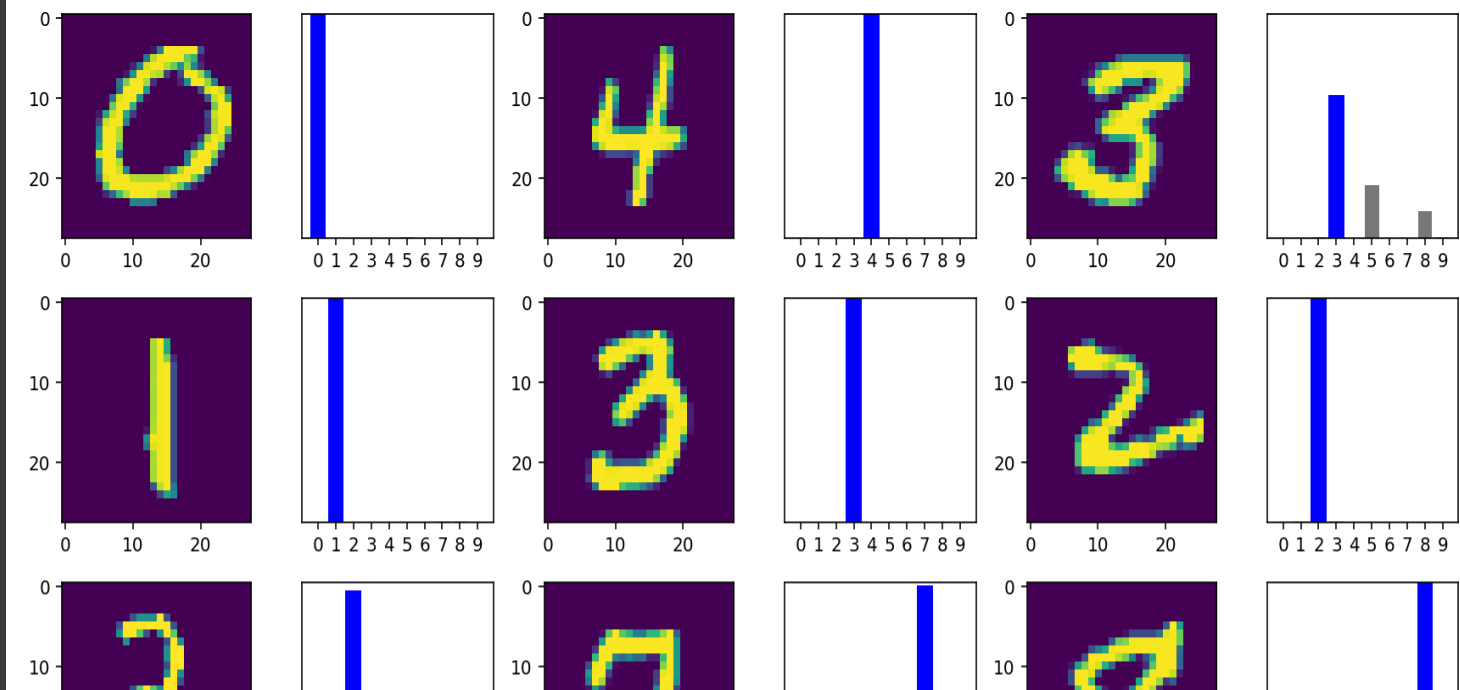
```

1
2 def plot_value_array(i, predictions_array, true_label):
3     predictions_array, true_label = predictions_array, true_label[i]
4     plt.grid(False)
5
6     plt.xticks(range(10))
7     plt.yticks([1])
8     thisplot = plt.bar(range(10), predictions_array, color="#777777")
9     plt.ylim([0, 1])
10    predicted_label = np.argmax(predictions_array)
11
12    thisplot[predicted_label].set_color('red')
13    thisplot[true_label].set_color('blue')
```

```

1  # Plot the first X test images, their predicted labels, and the true labels.
2  # Color correct predictions in blue and incorrect predictions in red.
3  num_rows = 5
4  num_cols = 3
5  num_images = num_rows*num_cols
6  plt.figure(figsize=(2*2*num_cols, 2*num_rows))
7  for i in range(num_images):
8      plt.subplot(num_rows, 2*num_cols, 2*i+1)
9      plt.imshow(x_test[i].reshape(28,28))
10     #plot_image(i, predictions[i], test_labels, test_images)
11     plt.subplot(num_rows, 2*num_cols, 2*i+2)
12     plot_value_array(i, predictions[i], y_test)
13 plt.tight_layout()
14 plt.show()
```





Observation

As we can see that the 3 is misclassified as 7 this may be because due to less training data . moreover that 3 looks much like 7.

All other test points are classified correctly



Comparison with PCA features

With only single layer



```

1  #PCA
2
3  data = pd.read_csv('17078PCA.csv')
4  pca_x = data.iloc[:,0:25].values
5  pca_y= data.iloc[:,25].values
6  print(len(pca_y))
7
8  #Splitting the data in training and test data
9  from sklearn.model_selection import train_test_split
10  pca_x_train,pca_x_test,pca_y_train,pca_y_test = train_test_split(pca_x,pca_y, test_size = 0.25, random_sta
11
12

```

3000

For this model, I have kept 25% of train data for testing and I has only single layer with input dimension as 25 and output is 10. I have used softmax as activation function.

```

1  import keras
2  from tensorflow.keras.models import Sequential
3  from tensorflow.keras.layers import Dense
4  from torch import nn
5  pca = Sequential()
6

```

```
6  
7  pca.add(Dense(units = 10 ,kernel_initializer = 'uniform',activation = 'softmax', input_dim =25))  
8
```


I am using stochastic gradient descent for optimizing with learning rate of 0.01 and
sparse_categorical_crossentropy as a loss function

```
1  pca.compile(optimizer=tf.keras.optimizers.Adam(0.01),loss='sparse_categorical_crossentropy',metrics=['accuracy'])  
2  history1 =[]  
3  history1.append(pca.fit(pca_x_train,pca_y_train,batch_size=64,epochs = 20,validation_data=(pca_x_test,pca_y_test)))  
4
```



```
Epoch 1/20
36/36 [=====] - 0s 5ms/step - loss: 1.0350 - accuracy: 0.7058 - val_loss: 0.5933 - v
Epoch 2/20
36/36 [=====] - 0s 3ms/step - loss: 0.5313 - accuracy: 0.8400 - val_loss: 0.5193 - v
Epoch 3/20
36/36 [=====] - 0s 3ms/step - loss: 0.4770 - accuracy: 0.8564 - val_loss: 0.4910 - v
Epoch 4/20
```

```
1 summarize_diagnostics(history1)
```



```
36/36 [=====] - 0s 3ms/step - loss: 0.4119 - accuracy: 0.8738 - val_loss: 0.4886 - v
```

▼ Observation

With only one layer the accuracy, and 100 epoch, the training accuracy converges to ~ 88% even the validation accuracy is around 83% which is much less than the SVM model that I have done in previous assignmnet. It had a test accuracy of 96% . Now lets try adding more layers Here I have made a fully connected network. The input and hidden layer has **sigmoid** activation function, whereas the output layer has **softmax** activation function.

```
36/36 [=====] - 0s 3ms/step - loss: 0.3982 - accuracy: 0.8769 - val_loss: 0.4889 - v
```

```
1 import keras
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense
4 from torch import nn
5 pca_with_hidden_layer = Sequential()
6
7 pca_with_hidden_layer.add(Dense(units = 20 ,kernel_initializer='uniform',activation = 'sigmoid' , input_dim =25))
8 pca_with_hidden_layer.add(Dense(units=15,kernel_initializer='uniform',activation='sigmoid'))
9 pca_with_hidden_layer.add(Dense(units=10,kernel_initializer='uniform',activation='softmax'))
10
11
```

```
0.75 | / | train
```

I am using stochastic gradient descent for optimizing with leaning rate of 0.01 and sparse_categorical_crossentropy as a loss function

```
1 pca_with_hidden_layer.compile(optimizer=tf.keras.optimizers.Adam(0.1),loss='sparse_categorical_crossentropy')
2 history2=[]
3 a=history2.append(pca_with_hidden_layer.fit(pca_x_train,pca_y_train,batch_size=64,epochs = 20,validation_data=(pca_x_val,pca_y_val)))
4
```




Epoch 1/20

36/36 [=====] - 0s 5ms/step - loss: 2.0839 - accuracy: 0.1916 - val_loss: 1.8033 - val_accuracy: 0.1916

Epoch 2/20

36/36 [=====] - 0s 4ms/step - loss: 1.5856 - accuracy: 0.3907 - val_loss: 1.2557 - val_accuracy: 0.3907

Epoch 3/20

36/36 [=====] - 0s 4ms/step - loss: 1.0475 - accuracy: 0.6329 - val_loss: 0.9052 - val_accuracy: 0.6329

Epoch 4/20

36/36 [=====] - 0s 4ms/step - loss: 0.7587 - accuracy: 0.7480 - val_loss: 0.7637 - val_accuracy: 0.7480

Epoch 5/20

36/36 [=====] - 0s 3ms/step - loss: 0.6174 - accuracy: 0.8031 - val_loss: 0.7239 - val_accuracy: 0.8031

Epoch 6/20

36/36 [=====] - 0s 3ms/step - loss: 0.5501 - accuracy: 0.8347 - val_loss: 0.6656 - val_accuracy: 0.8347

Epoch 7/20

36/36 [=====] - 0s 3ms/step - loss: 0.5362 - accuracy: 0.8262 - val_loss: 0.6900 - val_accuracy: 0.8262

Epoch 8/20

36/36 [=====] - 0s 3ms/step - loss: 0.5160 - accuracy: 0.8347 - val_loss: 0.6541 - val_accuracy: 0.8347

Epoch 9/20

36/36 [=====] - 0s 3ms/step - loss: 0.4604 - accuracy: 0.8547 - val_loss: 0.6686 - val_accuracy: 0.8547

Epoch 10/20

36/36 [=====] - 0s 4ms/step - loss: 0.4556 - accuracy: 0.8556 - val_loss: 0.6315 - val_accuracy: 0.8556

Epoch 11/20

36/36 [=====] - 0s 4ms/step - loss: 0.4491 - accuracy: 0.8556 - val_loss: 0.6257 - val_accuracy: 0.8556

Epoch 12/20

36/36 [=====] - 0s 3ms/step - loss: 0.4463 - accuracy: 0.8627 - val_loss: 0.6247 - val_accuracy: 0.8627

Epoch 13/20

36/36 [=====] - 0s 3ms/step - loss: 0.4567 - accuracy: 0.8484 - val_loss: 0.6527 - val_accuracy: 0.8484

Epoch 14/20

36/36 [=====] - 0s 3ms/step - loss: 0.4083 - accuracy: 0.8680 - val_loss: 0.6469 - val_accuracy: 0.8680

Epoch 15/20

36/36 [=====] - 0s 3ms/step - loss: 0.3937 - accuracy: 0.8738 - val_loss: 0.5943 - val_accuracy: 0.8738

Epoch 16/20

36/36 [=====] - 0s 3ms/step - loss: 0.3601 - accuracy: 0.8880 - val_loss: 0.6415 - val_accuracy: 0.8880

Epoch 17/20

36/36 [=====] - 0s 3ms/step - loss: 0.3711 - accuracy: 0.8876 - val_loss: 0.6130 - val_accuracy: 0.8876

Epoch 18/20

36/36 [=====] - 0s 3ms/step - loss: 0.3533 - accuracy: 0.8836 - val_loss: 0.7074 - val_accuracy: 0.8836


Epoch 19/20

36/36 [=====] - 0s 4ms/step - loss: 0.3950 - accuracy: 0.8698 - val_loss: 0.6776 - val_accuracy: 0.8698

Epoch 20/20

36/36 [=====] - 0s 3ms/step - loss: 0.3630 - accuracy: 0.8822 - val_loss: 0.6807 - val_accuracy: 0.8822

1 summarize_diagnostics(history2)



▼ Observation


After 20 epoch accuracy has increased as I have increased the depth of of the nural network. It is because by adding the hidden layer I gave the network flexibility to fit to the data but also increased the complexity of the basis function.

```
1 test_loss,test_acc = pca_with_hidden_layer.evaluate(pca_x_test,pca_y_test,verbose=2)
2 probability_model = tf.keras.Sequential([pca_with_hidden_layer])
3 predictions = probability_model.predict(pca_x_test)
```

24/24 - 0s - loss: 0.6953 - accuracy: 0.8240


Visualalization of the above model

Ploting the graph with true label, I havent shown the image based on the features given in the data set because they don't make sense to human. It is something like this Therefore I have printed the true lable against the predicted value.




```
1 # Plot the first X test images, their predicted labels, and the true labels.
2 # Color correct predictions in blue and incorrect predictions in red.
3
4
5 for i in range(0,10):
6
7     plot_value_array(i, predictions[i], pca_y_test)
8     print('-----')
9     print('Test(Actual) value = ',pca_y_test[i])
10
11
```



Test(Actual) value = 0




Test(Actual) value = 4




Test(Actual) value = 3




Test(Actual) value = 1



Test(Actual) value = 3



Test(Actual) value = 2



Test(Actual) value = 2



▼ CNN for MNIST Handwritten Digit Classification

As suggested, I am downloading the full mnist data from [here](#). I downloaded all data set then i decompresses them and then using idx2numpy library i converted them to the numpy array.

```
1 !wget http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
2 !wget http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
3 !wget http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
4 !wget http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
```



```
--2020-06-28 14:05:05-- http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Resolving yann.lecun.com (yann.lecun.com)... 104.28.6.204, 104.28.7.204, 172.67.171.76, ...
Connecting to yann.lecun.com (yann.lecun.com)|104.28.6.204|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 28881 (28K) [application/x-gzip]
Saving to: 'train-labels-idx1-ubyte.gz.1'

train-labels-idx1-u 100%[=====>] 28.20K --.-KB/s in 0.01s

2020-06-28 14:05:05 (2.04 MB/s) - 'train-labels-idx1-ubyte.gz.1' saved [28881/28881]

--2020-06-28 14:05:12-- http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Resolving yann.lecun.com (yann.lecun.com)... 172.67.171.76, 104.28.6.204, 104.28.7.204, ...
Connecting to yann.lecun.com (yann.lecun.com)|172.67.171.76|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 9912422 (9.5M) [application/x-gzip]
Saving to: 'train-images-idx3-ubyte.gz.1'

train-images-idx3-u 100%[=====>] 9.45M 60.6MB/s in 0.2s

2020-06-28 14:05:12 (60.6 MB/s) - 'train-images-idx3-ubyte.gz.1' saved [9912422/9912422]

--2020-06-28 14:05:19-- http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Resolving yann.lecun.com (yann.lecun.com)... 104.28.6.204, 104.28.7.204, 172.67.171.76, ...
Connecting to yann.lecun.com (yann.lecun.com)|104.28.6.204|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1648877 (1.6M) [application/x-gzip]
Saving to: 't10k-images-idx3-ubyte.gz.1'

t10k-images-idx3-ub 100%[=====>] 1.57M --.-KB/s in 0.1s

2020-06-28 14:05:19 (14.1 MB/s) - 't10k-images-idx3-ubyte.gz.1' saved [1648877/1648877]

--2020-06-28 14:05:24-- http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Resolving yann.lecun.com (yann.lecun.com)... 104.28.6.204, 104.28.7.204, 172.67.171.76, ...
Connecting to yann.lecun.com (yann.lecun.com)|104.28.6.204|:80... connected.
```

```
1 !pip install idx2numpy
```



```
Collecting idx2numpy
  Downloading https://files.pythonhosted.org/packages/23/6b/abab4652eb249f432c62431907c8de32bdcedb/
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from idx2numpy) (1.18.5)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from idx2numpy) (1.12.0)
Building wheels for collected packages: idx2numpy
  Building wheel for idx2numpy (setup.py) ... done
  Created wheel for idx2numpy: filename=idx2numpy-1.2.2-cp36-none-any.whl size=8032 sha256=1c88c3da2cfd
  Stored in directory: /root/.cache/pip/wheels/7a/b5/69/3e0757b3086607e95db70661798fdf98a77a0bb79c
Successfully built idx2numpy
Installing collected packages: idx2numpy
Successfully installed idx2numpy-1.2.2
```

▼ converting MNIST data to numpy array.

```
1 import idx2numpy
2 import numpy as np
3 train_images = idx2numpy.convert_from_file('/content/drive/My Drive/Assignment4/train-images-idx3-ubyte.gz')
4 test_images = idx2numpy.convert_from_file('/content/drive/My Drive/Assignment4/t10k-images-idx3-ubyte.gz')
5 train_label = idx2numpy.convert_from_file('/content/drive/My Drive/Assignment4/train-labels-idx1-ubyte.gz')
6 test_label = idx2numpy.convert_from_file('/content/drive/My Drive/Assignment4/t10k-labels-idx1-ubyte.gz')
```

```
1 from keras.datasets import mnist
```

```

2 from matplotlib import pyplot
3 # load dataset
4
5 # summarize loaded dataset
6 print('Train: X=%s, y=%s' % (train_images.shape, train_label.shape))
7 print('Test: X=%s, y=%s' % (test_images.shape, test_label.shape))
8 # plot first few images
9 for i in range(9):
10     # define subplot
11     pyplot.subplot(330 + 1 + i)
12     # plot raw pixel data
13     pyplot.imshow(train_images[i], cmap=pyplot.get_cmap('gray'))
14 # show the figure
15 pyplot.show()


```



```

Train: X=(60000, 28, 28), y=(60000,)
Test: X=(10000, 28, 28), y=(10000,)

```



we know that the images are all pre-aligned (e.g. each image only contains a hand-drawn digit), that the images all have the same square size of 28×28 pixels, and that the images are grayscale. load the images and reshape the data arrays to have a single color channel.

```

1 trainX = train_images.reshape((train_images.shape[0], 28, 28, 1))
2 testX = test_images.reshape((test_images.shape[0], 28, 28, 1))

```

Using a one hot encoding for the class element of each sample, transforming the integer into a 10 element binary vector with a 1 for the index of the class value, and 0 values for all other classes.

```

1 from keras.utils import to_categorical
2 trainY = to_categorical(train_label)
3 testY = to_categorical(test_label)

```

```

1 def load_dataset():
2
3     # reshape dataset to have a single channel
4     trainX = train_images.reshape((train_images.shape[0], 28, 28, 1))
5     testX = test_images.reshape((test_images.shape[0], 28, 28, 1))
6     # one hot encode target values
7     trainY = to_categorical(train_label)
8     testY = to_categorical(test_label)
9     return trainX, trainY, testX, testY

```

▼ Rescaling image value from [0,255] to [0,1]

```
1  # scale pixels
2  def prep_pixels(train, test):
3      # convert from integers to floats
4      train_norm = train.astype('float32')
5      test_norm = test.astype('float32')
6      # normalize to range 0-1
7      train_norm = train_norm / 255.0
8      test_norm = test_norm / 255.0
9      # return normalized images
10     return train_norm, test_norm
```

▼ Defining Model

For convolution front-end , I will start with single Convolution layer with small filter size (3,3) and a modest number of filter (32) followed by a max polling layer. The filter maps can then be flattened to provide features to the classifier . I will be using Softmax activation function. in between the feature extrator and output layer I will adda dense layer to interpret the features, 100 nodes

All layer will use reLU activation function and He weight initialization scheme.

I will use stochastic gradient decent optimizer with learning rate of 0.01 and momentum 0.9. The categorical cross-entropy loss function will be optimized, suitable for multi-class classification, and I will monitor the classification accuracy metric, which is appropriate given we have the same number of examples in each of the 10 classes.

```
1  # define cnn model
2  def define_model():
3      model = Sequential()
4      model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(28, 28, 1)))
5      model.add(MaxPooling2D((2, 2)))
6      model.add(Flatten())
7      model.add(Dense(100, activation='relu', kernel_initializer='he_uniform'))
8      model.add(Dense(10, activation='softmax'))
9      # compile model
10     opt = SGD(lr=0.01, momentum=0.9)
11     model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
12     return model
```

The evaluate_model() function below implements these behaviors, taking the training dataset as arguments and returning a list of accuracy scores

```
1
2  from numpy import mean
3  from numpy import std
4  from sklearn.model_selection import KFold
5  from sklearn.model_selection import KFold
6  from keras.datasets import mnist
7  from keras.utils import to_categorical
```



```

8 from keras.models import Sequential
9 from keras.layers import Conv2D
10 from keras.layers import MaxPooling2D
11 from keras.layers import Dense
12 from keras.layers import Flatten
13 from keras.optimizers import SGD
14 def evaluate_model(dataX, dataY, n_folds=5):
15     scores, histories = list(), list()
16     # prepare cross validation
17     kfold = KFold(n_folds, shuffle=True, random_state=1)
18     # enumerate splits
19     for train_ix, test_ix in kfold.split(dataX):
20         # define model
21         model = define_model()
22         # select rows for train and test
23         trainX, trainY, testX, testY = dataX[train_ix], dataY[train_ix], dataX[test_ix], dataY[test_ix]
24         # fit model
25         history = model.fit(trainX, trainY, epochs=10, batch_size=32, validation_data=(testX, testY), verbose=0)
26         # evaluate model
27         _, acc = model.evaluate(testX, testY, verbose=0)
28         print('> %.3f' % (acc * 100.0))
29         # stores scores
30         scores.append(acc)
31         histories.append(history)
32     return scores, histories

```

```

1 # summarize model performance
2 def summarize_performance(scores):
3     # print summary
4     print('Accuracy: mean=%.3f std=%.3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))
5     # box and whisker plots of results
6     pyplot.boxplot(scores)
7     pyplot.show()

```

```

1 # summarize model performance
2 def summarize_performance(scores):
3     # print summary
4     print('Accuracy: mean=%.3f std=%.3f, n=%d' % (mean(scores)*100, std(scores)*100, len(scores)))
5     # box and whisker plots of results
6     pyplot.boxplot(scores)
7     pyplot.show()
8

```

```

1
2 def run_test_harness():
3     # load dataset
4     trainX, trainY, testX, testY = load_dataset()
5     # prepare pixel data
6     trainX, testX = prep_pixels(trainX, testX)
7     # evaluate model
8     scores, histories = evaluate_model(trainX, trainY)
9     # learning curves
10    summarize_diagnostics(histories)
11    # summarize estimated performance
12    summarize_performance(scores)

```

```


1 run_test_harness()

```




```
> 98.442
> 98.667
> 98.608
> 98.850
> 98.742
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: MatplotlibDeprecationWarning: Adding an axe
if sys.path[0] == '':
```



Accuracy: mean=98.662 std=0.136, n=5



The model achieves a good fit, with train and test learning curves converging. There is no obvious sign of over- or underfitting. A box and whisker plot is created to summarize the distribution of accuracy scores.

Yes these are more natural or intuitive than the representations learnt by the standard neural net.

Batch normalization can be used after convolutional and fully connected layers. It has the effect of changing the distribution of the output of the layer, specifically by standardizing the outputs. This has the effect of stabilizing and accelerating the learning process.

As we increase the depth of the network the accuracy increases to 99%